



# An Inter-Cluster Communication Facility for Lightweight Manycore Processors in the Nanvix OS

João Vicente Souto

`joao.vicente.souto@grad.ufsc.br`

Graduação em Ciência da Computação

Depto. de Informática e Estatística

Universidade Federal de Santa Catarina - Florianópolis

Orientador: Prof. Márcio Bastos Castro, Dr.

Coorientador: Pedro Henrique Penna, Me.

# Presentation Outline

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

### 1 Introduction

### 2 Goals

### 3 Background

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

### 4 Development

- Low-Level Communication
- User-Level Communication

### 5 Experiments

### 6 Conclusions

**Introduction**

Goals

**Background**

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

**Development**

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

Introduction

# Historical Evolution and Trend of Processors

## Introduction

## Goals

## Background


Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions



Increase in **processor performance**

# Historical Evolution and Trend of Processors

## Introduction

## Goals

## Background


Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions



Limited by **frequency barrier**

# Historical Evolution and Trend of Processors

## Introduction

## Goals

## Background


Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions



Softened by **technological advancement**

# Historical Evolution and Trend of Processors

## Introduction

## Goals

## Background


Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions



Emergence of **multicores** and **manycores**

# Historical Evolution and Trend of Processors

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



Advent of **lightweight manycores**



# Lightweight Manycores Particularities

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



Overview of a Manycore

- **Hundreds of Lightweight Cores**
  - Expose Massive thread-level parallelism
  - Feature low-power consumption
  - Target MIMD workloads
- Distributed Memory Architecture
- On-Chip Heterogeneity

# Lightweight Manycores Particularities

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



Overview of a Manycore

- Hundreds of Lightweight Cores
- **Distributed Memory Architecture**
  - Grants scalability
  - Relies on a Network-on-Chip (NoC)
  - Has constrained memory systems
- On-Chip Heterogeneity

# Lightweight Manycores Particularities

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



## Overview of a Manycore

- Hundreds of Lightweight Cores
- Distributed Memory Architecture
- **On-Chip Heterogeneity**
  - Features different components



Introduction

Goals

Background

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

Development

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

Goals

# Context and Goals

## Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development

Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

## ■ Work context

- Challenges arising from **runtimes** and **OSs**
- OSs for **next-generation** for lightweight manycores
- **Nanvix OS**

## ■ Goals

- Definition and proposal of an **Inter-Cluster Communication Interface** for lightweight manycores
- Implementation in the **Nanvix HAL** on Kalray MPPA-256 Lightweight Manycore Processor
- Integration with the **Nanvix Microkernel**
- Performance evaluation using synthetic micro-benchmarks with **collective communication routines**



Introduction

Goals

**Background**

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

Development

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

# Background

# Models of Operating Systems (OS) for Multicores

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Replicated OS
- Master-Slave OS
- Symmetric OS



Each core has a **copy of the OS**

**Less concurrency** but needs a lot of memory

# Models of Operating Systems (OS) for Multicores

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Replicated OS
- **Master-Slave OS**
- Symmetric OS



**Asymmetric execution** where only one core handles the OS

**Better scalability** and **less inter-core interference** but inserts a possible **bottleneck**



# Models of Operating Systems (OS) for Multicores

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Replicated OS
- Master-Slave OS
- **Symmetric OS**



**OS shared** by all cores

**Efficient with few cores** and **cache consistency** but has **inter-core interference**

# Low-Level Communication for Message-Passing Model

## Introduction

## Goals

## Background

Operating Systems Models

**Message-Passing Model**

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Computer network concepts**
- Network interfaces
- Performance impacts
  - Number of intermediate copies
  - Direct Memory Access (DMA)

# User-Level Communication for Message-Passing Model

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication


## Experiments

## Conclusions

### ■ Send/receive primitives

#### ■ Synchronous calls

#### ■ Asynchronous calls



**Requester waits** for task to be completed

# User-Level Communication for Message-Passing Model

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Send/receive primitives
  - Synchronous calls
  - **Asynchronous calls**



**Requester is released** to continue to run in parallel

# Kalray MPPA-256

## A Lightweight Manycore Processor

### Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development

Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

## ■ 288 processing cores


- 16 Compute Cluster (CC)
- 4 I/O Cluster (IO)

## ■ Data NoC (D-NoC)

- 256 RX slots
- 8 TX channels
- 8  $\mu$ threads for async TX

## ■ Control NoC (C-NoC)

- 128 RX slots
- 4 TX channels



# The Nanvix Operating System

## Overview - The Nanvix Project

### Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development

Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

## ■ Open source and collaborative project

- Home-grown instructional OS
- 4 Professors (Brazil and France)
- 1 PhD, 1 MSc and 3 BSc Students
- 9 past contributors
- UGA, PUC Minas, UFSC, and Grenoble INP




*Bingo, our mascot*



## ■ Long-Term Goal

- General purpose OS
- POSIX-Compliant



# Nanvix Multikernel

## Multikernel OS Structure

### Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development


Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

- **OS services are system processes** and run in isolation
- User processes **request services via message-passing**



# Nanvix Microkernel

## Asymmetric Microkernel

### Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development


Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

- Follows a **Master-Slave OS** model
- Provides **bare-bones of system abstractions**
- Rich system call interface (Kernel Call)





# Nanvix Hardware Abstraction Layer (HAL)

Kernel portability

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

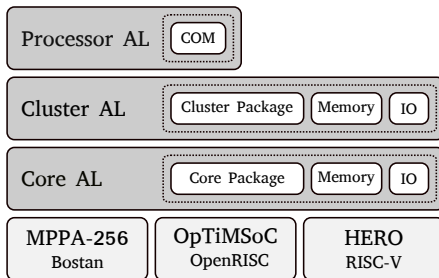
Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Generic and flexible** hardware abstraction layer
- **Standard view** of these emerging processors



Introduction

Goals

Background

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

Development

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

Development

# Inter-Cluster Communication Interface

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

### ■ Three communication abstractions

- Sync
  - Mailbox
  - Portal
- 
- More precise
  - Easy-to-use
  - Scalable
  - Easily portable

# MPPA-256 Hardware Features

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Hardware independence** in resource identification
  - NoC interfaces (Physical to Logical)
  - NoC resources (Partitioned by abstraction)
- Low-level communication depends on **two features**
  - Interrupt System
  - DMA
- **Work limitations**
  - Poor documentation and examples
  - Reverse engineering
  - Hardware bugs
  - **No asynchronous sends**

# General Concepts of Comm. Abstractions

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Naming Convention** (suffixes)
  - Sender: open/close/awrite/signal/wait.
  - Receiver: create/unlink/aread/wait
- Interfaces export **only asynchronous calls**
- Master core **cannot be blocked**



**Lazy transfer behavior**

# General Concepts of Comm. Abstractions

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Naming Convention** (suffixes)
  - Sender: open/close/awrite/signal/wait.
  - Receiver: create/unlink/aread/wait
- Interfaces export **only asynchronous calls**
- Master core **cannot be blocked**



**Master** without sending permission **keeps the parameters**

# General Concepts of Comm. Abstractions

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Naming Convention** (suffixes)
  - Sender: open/close/awrite/signal/wait.
  - Receiver: create/unlink/aread/wait
- Interfaces export **only asynchronous calls**
- Master core **cannot be blocked**



**Master sends data upon allow and releases slave**

# Sync Abstraction

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Provides cluster synchronization across **distributed barriers**
- Similarly to **POSIX Signals**
- Two modes
  - **ALL\_TO\_ONE**
  - **ONE\_TO\_ALL**





# Sync Abstraction

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Provides cluster synchronization across **distributed barriers**
- Similarly to **POSIX Signals**
- Two modes
  - **ALL\_TO\_ONE**
  - **ONE\_TO\_ALL**



# Sync Abstraction Implementation

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Uses only **C-NoC** Resources
- **Node list** of involved Logic IDs
- Master node must always be the first
- RX resources **related to the Master ID**

# Mailbox Abstraction

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Allows exchange **fixed-length message**
- Similarly to **POSIX Message Queue**
- Receiver allocates space to N messages
- Sender transfer to predefined location



# Mailbox Abstraction Implementation

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Merges the use of **D-NoC** and **C-NoC**
- Message queue allocated within the kernel space
- **Quality of Service**
  - **One message** per NoC Node



# Portal Abstraction

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Allows exchange **arbitrary amounts of data**
- Similarly to **POSIX Pipes**
- **One-way** channel for data transfer



# Portal Abstraction Implementation

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Merges the use of **D-NoC** and **C-NoC**
- No intermediate copies
- **Quality of Service**
  - *Receiver notifies sender that it is able to receive*



# Integration with Nanvix Microkernel

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Rich resource management** is forwarded to the microkernel
- Nanvix Microkernel seek **to provide**:
  - Protection
  - Management
  - Multiplexing
- **Impacts of the Master-Slave OS Model**
  - Only master changes the internal structures of the OS
  - Slaves allowed to read lock addresses to wait for async calls

# Microkernel

## Protection, Management and Multiplexing

### Introduction

### Goals

### Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

### Development

Low-Level Communication

User-Level Communication

### Experiments

### Conclusions

## ■ Protection and management involve two phases

### ■ Slave phase

- valid file descriptors
- non-null buffer pointers
- buffer sizes within the stipulated limit

### ■ Master phase

- conflicting operations
- measures communication parameters
- interacts with Nanvix HAL detecting errors

## ■ Multiplexing

- Identifies creations/openings with **same arguments**
- Structures keep a **reference counter**
- Operations of read/write set resources **to busy**
- Forces **serialization** of the operations



# Validation and Correctness Tests

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Extra effort in **ensuring semantics** of implementations **across multiple architectures**
- **API tests**
  - Arguments within valid value ranges
  - Operations with valid semantic
- **FAULT tests**
  - Arguments outside domain range
  - Operations with invalid semantic
  - Expected error value



Introduction

Goals

Background

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

Development

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

# Experiments

## Introduction

## Goals

## Background

Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS


## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions

- Evaluation of the performance of **data transfer services**
  - Sync + Mailbox
  - Sync + Portal
- Four well-known **collective communication routines**
  - **Broadcast**
  - Gather
  - AllGather
  - Ping-Pong



## Introduction

## Goals

## Background

Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS


## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions

- Evaluation of the performance of **data transfer services**
  - Sync + Mailbox
  - Sync + Portal
- Four well-known **collective communication routines**
  - Broadcast
  - **Gather**
  - AllGather
  - Ping-Pong



## Introduction

## Goals

## Background

Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS


## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions

- Evaluation of the performance of **data transfer services**
  - Sync + Mailbox
  - Sync + Portal
- Four well-known **collective communication routines**
  - Broadcast
  - Gather
  - **AllGather**
  - Ping-Pong



## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- Evaluation of the performance of **data transfer services**
  - Sync + Mailbox
  - Sync + Portal
- Four well-known **collective communication routines**
  - Broadcast
  - Gather
  - AllGather
  - Ping-Pong



# Experimental Design

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

- **Throughput of the Portal**
  - Constant: 1 IO and 16 CCs
  - Variable: 4, 8, 16, 32, 64 KB
- **Latency of the Mailbox**
  - Constant: 120 B
  - Variable: 1 IO and 1 to 16 CCs
- **50 iterations**, first 10 discarded to warmup
- Metrics do not represent an aggregation
- Standard error inferior to **1%**

# Portal Analysis

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



**Three distinct behaviors**



# Portal Analysis

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



**Optimum size** between 8 KB and 16 KB

# Mailbox Analysis

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



**Three distinct behaviors too**

# Mailbox Analysis

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



Well-known **distributed algorithms can be efficiently supported** by Nanvix OS



Introduction

Goals

Background

- Operating Systems Models
- Message-Passing Model
- Kalray MPPA-256
- Nanvix OS

Development

- Low-Level Communication
- User-Level Communication

Experiments

Conclusions

# Conclusions

# Conclusions

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

### ■ Motivation

- Historical evolution from single-cores to Lightweight Manycores

### ■ Contribution


- A Inter-Cluster Communication Facility for LW Processors

### ■ Results

- **Optimal sizes** for large data transfers
- Well-known **distributed algorithms** can be efficiently **supported** by Nanvix OS

### ■ Future Works on Nanvix OS

- Remove limitation on asynchronous send
- MPI port (BSc dissertation)
- Shared Memory Service (MSc dissertation)
- **Distributed Process Scheduling (MSc dissertation)**



Thank you!  
Questions?

João Vicente Souto

[joao.vicente.souto@grad.ufsc.br](mailto:joao.vicente.souto@grad.ufsc.br)

Graduação em Ciência da Computação  
Depto. de Informática e Estatística

Universidade Federal de Santa Catarina - Florianópolis

Orientador: Prof. Márcio Bastos Castro, Dr.  
Coorientador: Pedro Henrique Penna, Me.

# References I

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development


Low-Level Communication

User-Level Communication


## Experiments

## Conclusions

 BAUMANN, A. et al. The Multikernel: A New OS Architecture for Scalable Multicore Systems. In: *Proceedings of the 22nd ACM Symposium on Operating Systems Principles*. Big Sky, Montana, USA: ACM, 2009. (SOSP '09), p. 29–44. ISBN 978-1-60558-752-3. Disponível em: <https://dl.acm.org/citation.cfm?doid=1629575.1629579>.

 DINECHIN, B. D. de et al. A Distributed Run-Time Environment for the Kalray MPPA-256 Integrated Manycore Processor. In: *Procedia Computer Science*. Barcelona, Spain: Elsevier, 2013. (ICCS '13, v. 18), p. 1654–1663. ISBN 1877-0509. Disponível em: <http://linkinghub.elsevier.com/retrieve/pii/S1877050913004766>.

 KENDALL, W.; NATH, D.; BLAND, W. *A Comprehensive MPI Tutorial Resource*. 2019. <https://mpitutorial.com/>, Last accessed on 2019-10-22.

 PENNA, P. H. et al. Using the Nanvix Operating System in Undergraduate Operating System Courses. In: *2017 VII Brazilian Symposium on Computing Systems Engineering*. Curitiba, Brazil: IEEE, 2017. (SBESC '17), p. 193–198. ISBN 978-1-5386-3590-2. Disponível em: <http://ieeexplore.ieee.org/document/8116579/>.

# References II

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development


Low-Level Communication


User-Level Communication

## Experiments


## Conclusions

 PENNA, P. H. et al. Using The Nanvix Operating System in Undergraduate Operating System Courses. In: *VII Brazilian Symposium on Computing Systems Engineering*. Curitiba, Brazil: [s.n.], 2017. Disponível em: <https://hal.archives-ouvertes.fr/hal-01635880>.

 PENNA, P. H. et al. RMem: An OS Service for Transparent Remote Memory Access in Lightweight Manycores. In: *MultiProg 2019 - 25th International Workshop on Programmability and Architectures for Heterogeneous Multicores*. Valencia, Spain: [s.n.], 2019. (High-Performance and Embedded Architectures and Compilers Workshops (HiPEAC Workshops)), p. 1–16. Disponível em: <https://hal.archives-ouvertes.fr/hal-01986366>.

 PENNA, P. H. et al. An os service for transparent remote memory accesses in noc-based lightweight manycores. Poster. 2018.

 RUPP, K. *Microprocessor Trend Data*. 2018. <https://github.com/karlrupp/microprocessor-trend-data>, Last accessed on 2019-06-26.

 SILBERSCHATZ, A.; GALVIN, P. B.; GAGNE, G. *Operating System Concepts*. 9th. ed. [S.l.]: Wiley Publishing, 2012. ISBN 1118063333, 9781118063330.



# References III

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development


Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

 TANENBAUM, A. S.; BOS, H. *Modern Operating Systems*. 4th. ed. Upper Saddle River, NJ, USA: Prentice Hall Press, 2014. ISBN 013359162X, 9780133591620.

 WICKRAMASINGHE, U.; LUMSDAINE, A. A survey of methods for collective communication optimization and tuning. *CoRR*, abs/1611.06334, 2016. Disponível em: <http://arxiv.org/abs/1611.06334>.

# MIMD workload

## Introduction

## Goals

## Background


Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions



# Type of network-on-chip of the MPPA-256

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



Torus network

# Execution example of the Nanvix Microkernel

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



# POSIX Compliance example of Nanvix Multikernel

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS


## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions



# Network-on-Chip Identifiers

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

	Physical ID	Logical ID
IO 0	128-131	0-3
IO 1	192-195	4-7
CCs	0-15	8-23

# Resource Identifiers

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

	C-NoC		D-NoC	
	RX ID	TX ID	RX ID	TX ID
<b>Mailbox</b>	0-23	0	0-23	1-3
<b>Portal</b>	24-47	1-2	24-47	4-7
<b>Sync</b>	48-71	3	-	-

# Simplified NoC handler algorithm

## Introduction

## Goals

## Background

Operating Systems Models

Message-Passing Model

Kalray MPPA-256

Nanvix OS

## Development

Low-Level Communication

User-Level Communication

## Experiments

## Conclusions

**Require:**  $status[M_{Interfaces}][N_{Resources}]$ , interrupt status of a resource.

**Require:**  $handlers[M_{Interfaces}][N_{Resources}]$ , interrupt handler of a resource.

```
1: procedure NOC__HANDLER
2:   for  $i \in [1, M_{Interfaces}]$  do
3:     for  $j \in [1, N_{Resources}]$  do
4:       if  $status[i][j] == InterruptTriggered$  then
5:         CLEAN_STATUS( $i, j$ )
6:         HANDLERS[I][J]( $i, j$ )
```



# Simplified lazy transfer algorithm

## Introduction

## Goals

## Background

Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions

**Require:** *resources*, Abstraction Resource Table

Configures data transfer.

```
1: procedure ASYNC_WRITE(id, message, size)
2:   resources[id].message ← message
3:   resources[id].size ← size
4:   if resources[id].has_permission then
5:     DO_LAZY_TRANSFER(id)
6:   else
7:     resources[id].is_waiting ← True
```

Receives permission.

```
8: procedure ABSTRACTION_HANDLER(id)
9:   if resources[id].is_waiting then
10:    DO_LAZY_TRANSFER(id)
11:   else
12:    resources[id].has_permission ← True
```

Transfers the data.

```
13: procedure DO_LAZY_TRANSFER(id)
14:   resources[id].is_waiting ← False
15:   resources[id].has_permission ← False
16:   TRANSFER_DATA(resources[id].message, resources[id].size)
17:   UNLOCK(resources[id].lock)
```

▷ Releases slave core.

# References used

## Introduction

## Goals

## Background

Operating Systems Models  
Message-Passing Model  
Kalray MPPA-256  
Nanvix OS

## Development

Low-Level Communication  
User-Level Communication

## Experiments

## Conclusions

- Processor trend ([RUPP, 2018](#))
- Kalray MPPA-256 ([DINECHIN et al., 2013](#))
- Multikernel ([BAUMANN et al., 2009](#))
- Nanvix Project ([PENNA et al., 2017b](#); [PENNA et al., 2017a](#); [PENNA et al., 2019](#); [PENNA et al., 2018](#))
- Concepts about OSs ([TANENBAUM; BOS, 2014](#))
- Models of OSs ([SILBERSCHATZ; GALVIN; GAGNE, 2012](#))
- MPI collective communication routines ([WICKRAMASINGHE; LUMSDAINE, 2016](#); [KENDALL; NATH; BLAND, 2019](#))